

Simon Zambrovski & Jan Galinski - Holisticon AG



Meet AVRO!

The future of serialization for CQRS/ES systems using Axon Framework.

Who are we?

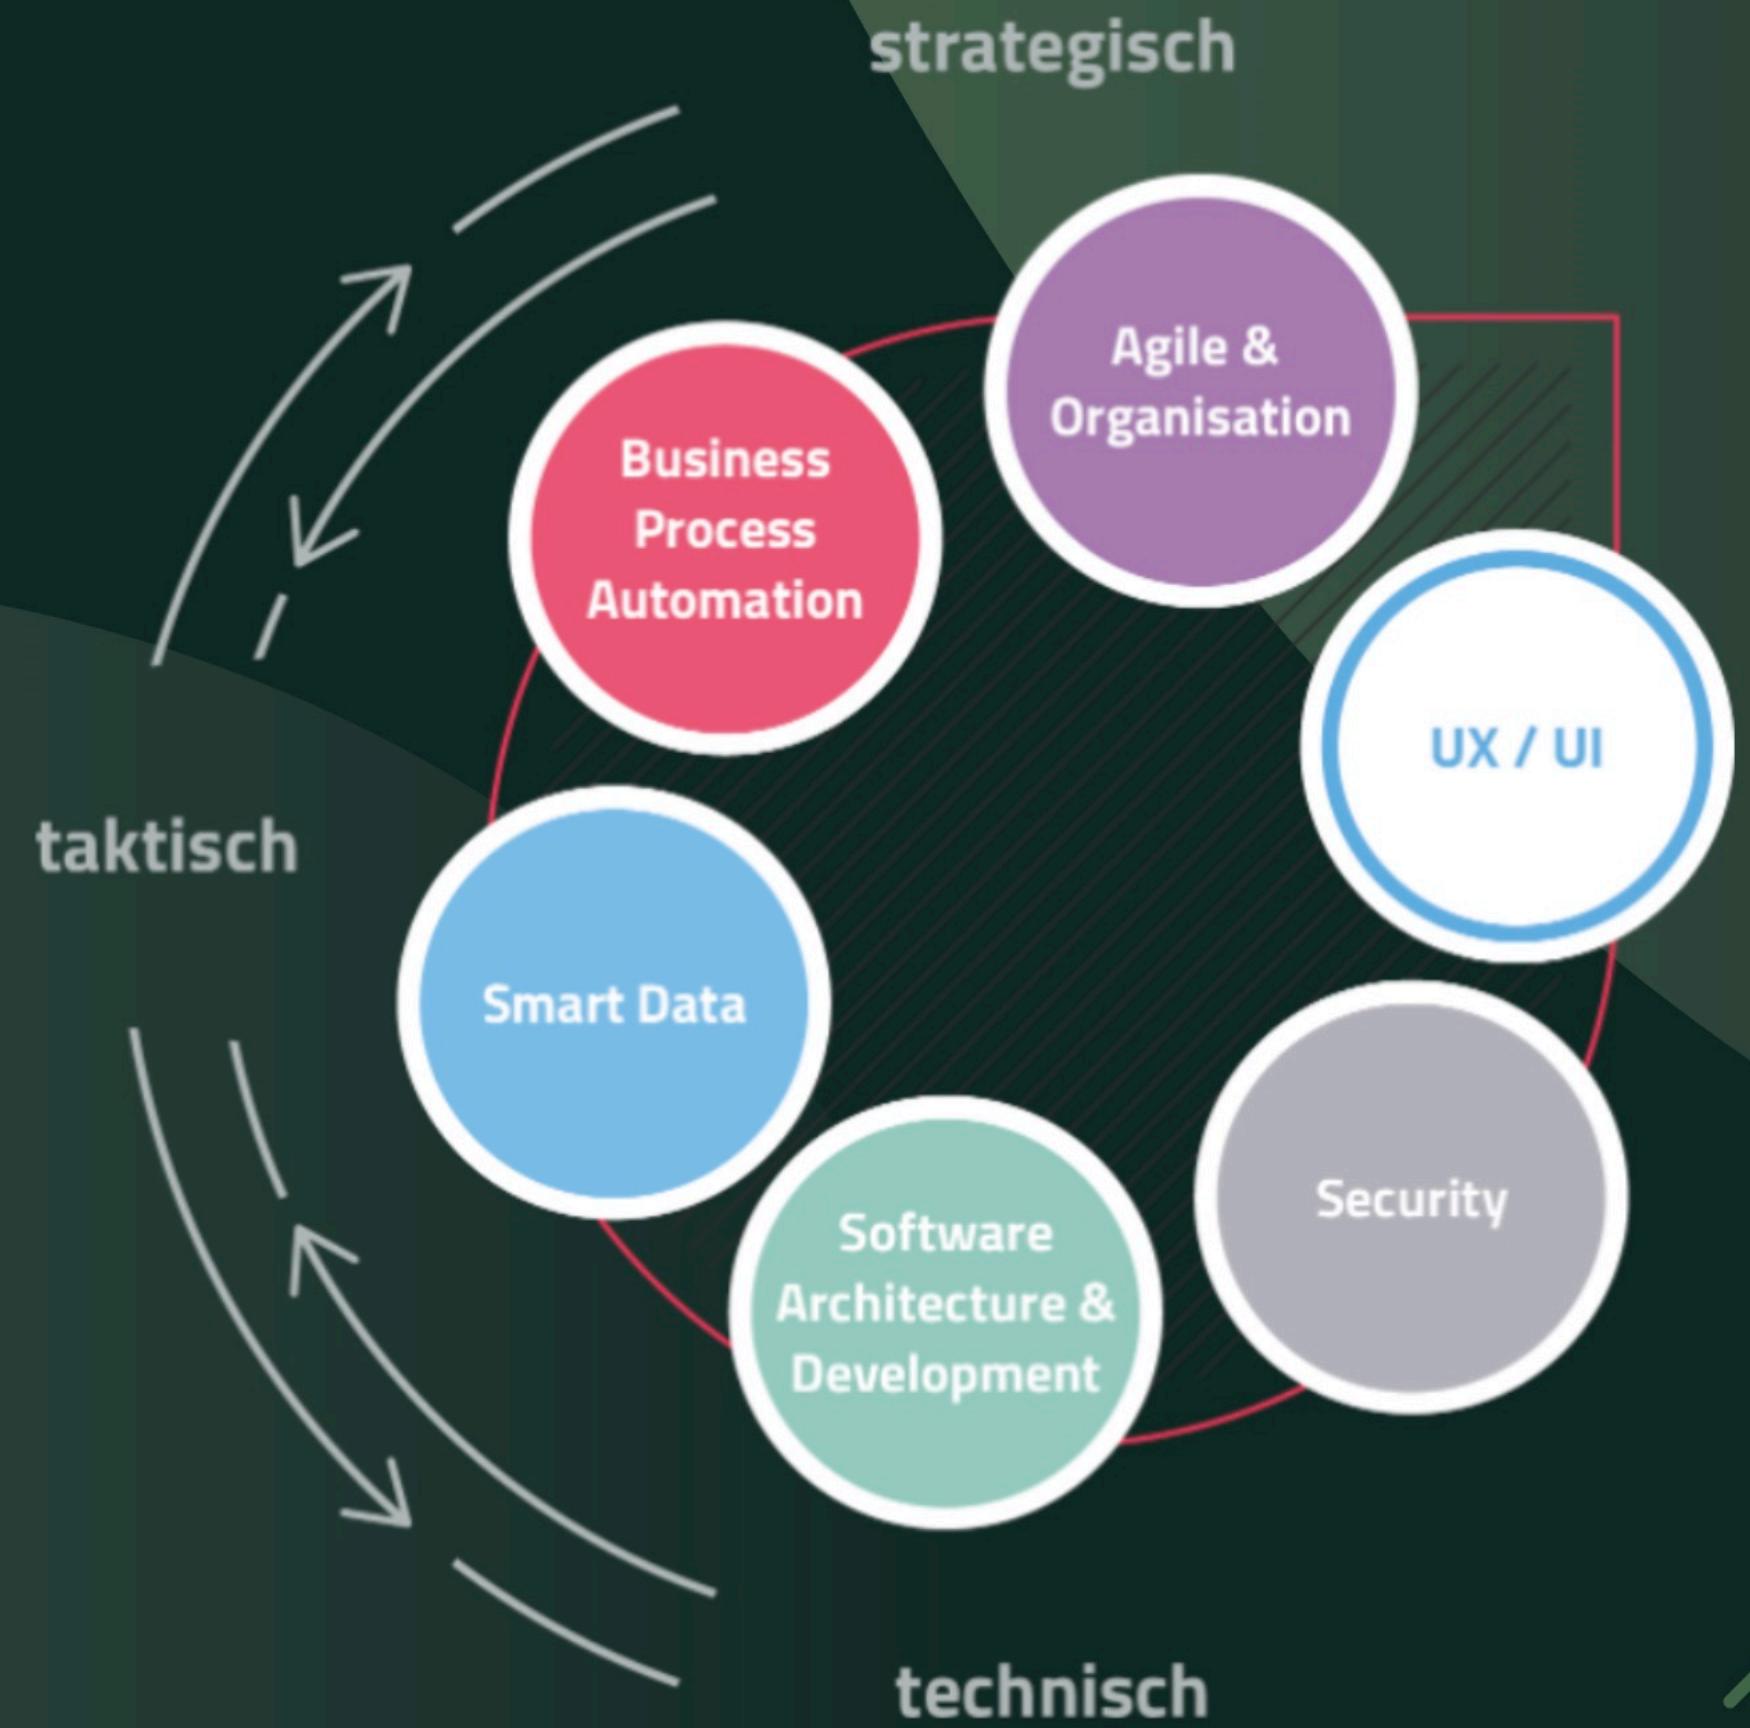


- Consultants @ holisticon
- Backend developers, software architects
- DDD, CQRS/ES, Axon Framework, Spring, Kotlin, Java



About Holisticon

- IT & Management Consulting
 - Hamburg
 - Hannover
 - Lissabon
- Founded 2007
- AxonIQ Partner
- Cool crew (~100 people)
- Kununu Best Employer 2023 in Germany
- And yes ...
 - we are hiring



DDD & CQRS/ES

The Challenges

Business Challenges

Design

- Messages are **first class citizens** (*Event Storming*)
- Your system **behavior** and state is expressed by commands, events and queries
 - Messages have **semantics** (a command triggers an event, a response answers a query)
- Business and IT need to **communicate** in terms of messages and **UBL**
 - Your Business does not work with String and Integer, it uses CustomerId and Amount

Evolution

- You will need to **identify** and read events years from now
 - Events are stored **long term** in an append only store
- If you are forced to change an event you need to know if it is still **compatible** to the previous version
- ...
- You do not want to rely on some developer **hacking** everything correctly in some IDE

IT Challenges

Development

- Your system will consist of multiple **micro-services** (*polyglot?*)
- You will need to **share** your commands, events and queries between services and teams.
 - You must always use the latest **revisions**
 - You must share **serialization** details as well
- The semantics of your messages and the **axon framework** require you to implement specific handler functions
- You must not make any **mistakes!**

Operation

- No overhead: high **performance** serialization
- Long term storage: You want to have a **compact storage format** that does not waste precious disk space
- If business changes events, you will have to develop smart **upcasting**, keeping compatibility
 - *Upcasters need to be distributed as well*

QA

- You need to **test** handlers, messages and upcasters

DDD & CQRS/ES

The Solution

What is Apache AVRO?

Self Promotion

Apache Avro™ is the leading serialization format for record data, and first choice for streaming data pipelines.

It offers excellent schema evolution, and has implementations for the JVM (Java, Kotlin, Scala, ...), Python, C/C++/C#, PHP, Ruby, Rust, JavaScript, and even Perl.

Fact sheet

- Created in 2011, as part of Hadoop ecosystem
- Platform independent - FE/BE, LLM, Messaging, ...
- Schema-First approach for types and protocols
 - Defines a type system (*int, string, long, array, record, ...*)
 - Defines serialization encodings (*JSON, Binary, Single Object Encoding*)

What is Apache AVRO?

Example of a simplified bank-account context (json):

```
235 "messages": {  
236     "createBankAccount": {  
237         "doc": "Creates a new bank account.",  
238         "meta": {  
239             "group": "BankAccountAggregate",  
240             "type": "deciderInit"  
241         },  
242         "request": [  
243             {  
244                 "name": "command",  
245                 "type": "CreateBankAccountCommand"  
246             }  
247         ],  
248         "response": "BankAccountCreatedEvent",  
249     }
```

How can Apache AVRO help you?

Design

- **Schema-First Approach**
 - Directly express your commands, events and queries via schema declaration
- **Describe context behavior as Protocols**
 - simple messages abstraction for request, response and error
- **centralized Schema-Registry**
- **UBL supported by Logical Types**
 - define custom value types (CustomerId, ...)
- **Extensible schema declaration**
 - Express (non)technical properties inside the schema

Evolution

- **Schema compatibility matrix**
 - different levels of forward/backward compatibility
 - separation of writer and reader schema
- **Schema fingerprint embedded in messages**
 - ability to look up the writer-schema
- **Simplified upcasting**
 - definition of **default** values for new properties
 - **Conflicts** can be determined on schema level
 - The **intermediate representation (*GenericRecord*)** understands the underlying schema (no mental model mismatch)

How can our AXON–AVRO help?

Development

- Do not share versioned **libraries** between teams, use **schema registry**
- Generate **Core API**
 - (data) classes for commands, events, queries, responses and errors
- Generate **Command Model**
 - interfaces for CommandHandler and EventSourcingHandler functions
 - custom CommandGateway for type safe publishing
- Generate **Query Model**
 - interfaces for QueryHandler functions
 - QueryGateway extensions for type safe and reliant query/response clients

Operation

- Pluggable AVRO serializer
 - auto configurable
 - about as fast as jackson
- **Compact storage via Single Object Encoded bytes**
 - reduce disk usage by min 10% < json
- Axon dashboard **plugin** to read events as json

QA

- Avoid making any **mistakes** because boilerplate is derived from **schema!**

Demo Time

aka video time

TODO



Summary

Summary

What we showed today

- Apache **AVRO** is a powerful schema first **serialization** platform
- Schema-First is a good idea for DDD/UBL and **Business/IT-Alignment**
- **Code-Generation** improves your life by avoiding boilerplate and bugs
- **Prototype** implementation of axon-avro - needs to improve

What we do next

- Implement missing **MUST features**
 - custom logical (value) types
 - registry support
- Explore/implement optional ideas
 - **jMolecules** support - annotated architecture
 - derive **documentation** (context map, plant-uml, markdown, ...)
 - combine with **FModel** - generate delegating aggregate
- Research
 - use it in polyglot environments
 - Data protection and schema?

Thank you!

Links



<https://github.com/holixon/axon-avro>

<https://avro.apache.org/>

Questions?



<https://holisticon.de/> – <https://about.me/zambrovski> – <https://about.me/jangalinski>